

Running Time and Growth

The *time complexity* $T(n)$ of an algorithm is the maximum amount of time taken on any input of size n .

$T(n) \in \mathcal{O}(g(n))$ if $\exists c > 0$ and $n_0 \geq 0$ s.t. $\forall n \geq n_0 \Rightarrow T(n) \leq c \cdot g(n)$.

$T(n) \in \Omega(g(n))$ if $\exists c > 0$ and $n_0 \geq 0$ s.t. $\forall n \geq n_0 \Rightarrow T(n) \geq c \cdot g(n)$.

$T(n) \in \Theta(g(n))$ if $T(n) \in \mathcal{O}(g(n))$ and $T(n) \in \Omega(g(n))$.

Order: $\mathcal{O}(\log_a n) < \mathcal{O}(n^d) < \mathcal{O}(r^n)$

$\mathcal{O}(\log n)$	binary search s. list
$\mathcal{O}(n)$	max element in list
$\mathcal{O}(n \log n)$	sort elements
$\mathcal{O}(n^2)$	find pair of closest
$\mathcal{O}(n^3)$	find disjoint sets
$\mathcal{O}(n^k)$	check $\exists e = (u, v)$
$\mathcal{O}(2^n)$	enum all subsets
$\mathcal{O}(n!)$	naive TSP

Graph Representation

Adjacency matrix: $n \times n$ matrix with $A_{uv} = 1$ if $(u, v) \in E$. Takes space n^2 . If $(u, v) \in E$ takes $\Theta(1)$ time.

Identifying all $e \in E$ takes $\Theta(n^2)$ time. *Adjacency list*: Node indexed array of lists. Takes space $m + n$. If $(u, v) \in E$ takes $\mathcal{O}(\deg(u))$ time. Identifying all $e \in E$ takes $\Theta(m + n)$ time.

Greedy Algorithms

Greedy-choice property: globally optimal solution with locally optimal choices.

Interval Scheduling

Goal: Find maximum subset of mutually compatible jobs.

```

1: procedure INTERVALSCHEDULING
2:   Sort by finish time
3:    $A \leftarrow \emptyset$ 
4:    $f_{j^*} = 0$ 
5:   for  $i = 1$  to  $n$  do
6:     if  $s_j \leq f_{j^*}$  then
7:        $A \leftarrow A \cup \{j\}$ 
8:        $f_{j^*} = f_j$ 
9:   return  $A$   $\triangleright \mathcal{O}(n \log n)$ 

```

Proof. Let $A : a_1, \dots, a_k$ be the greedy solution and let $O : o_1, \dots, o_m$ with $m \geq k$ be the optimal solution. We claim that $\forall r \leq k \Rightarrow f(a_r) \leq f(o_r)$. For $r = 1$, $f(a_1) \leq f(o_1)$ by greedy choice. Suppose

this holds for $r - 1 : f(a_{r-1}) \leq f(o_{r-1})$. Then $f(o_{r-1}) \leq s(o_r) \leq f(o_r)$ by solution feasibility. Thus $f(a_{r-1}) \leq s(o_r)$, so o_r was an option for greedy not picked. Since a_r was chosen instead, $f(a_r) \leq f(o_r)$. Now assume $k < m$ for contradiction. Then $f(a_k) \leq f(o_k) \leq s(o_{k+1})$ because O is feasible, so o_{k+1} is an option for A , which contradicts $|A| = k$. \square

Interval Partiton

Goal: Find minimum number of rooms to schedule all jobs without overlaps within a room.

```

1: procedure INTERVALPARTITION
2:   Sort by start time
3:    $d \leftarrow 0$ 
4:   for  $j = 1$  to  $n$  do
5:     if  $j$  compatible with class  $k$  then
6:       Add  $j$  to  $k$ 
7:     else
8:       Add new class room  $d + 1$ 
9:       Add  $j$  to  $d + 1$ 
10:       $d \leftarrow d + 1$   $\triangleright \mathcal{O}(n \log n)$ 

```

Proof. Let d be the number of classrooms from greedy. The last room d is opened because job j was incompatible with $d - 1$ rooms. These rooms have jobs ending after s_j . Thus we have d lectures overlapping at time $s_j + \epsilon$ and the depth is $\geq d$. But the number of rooms \geq depth, so greedy must have $d \geq$ depth, so solution with $d =$ depth is optimal. \square

Minimize Lateness

Goal: Schedule all jobs to minimize maximum lateness $L = \max \ell_j$.

```

1: procedure MINLATE
2:   Sort by increasing deadline
3:    $t \leftarrow 0$ 
4:   for  $i = j$  to  $n$  do
5:     Assign  $j$  to  $[t, t + t_j]$ 
6:      $s_j \leftarrow t$ ,  $f_j \leftarrow t + t_j$ 
7:      $t \leftarrow t + t_j$ 
8:   return Intervals  $[s_j, f_j]$   $\triangleright \mathcal{O}(n \log n)$ 

```

Def. An *inversion* is a pair of jobs i and j with $d_i < d_j$ but j scheduled before i .

Proof. Note that greedy has no idle time, $\exists O$ with no idle time, and greedy has no inversions. All schedules with no inversions and no idle time have the same lateness, because all jobs with same d come in a block of consecutive jobs if no inversions. Any reordering of these jobs with retain the last job having the lateness $f - d$. Let ℓ be the lateness before the swap and ℓ' be it after. $\ell'_k = \ell_k \forall k \neq i, j$. $\ell'_i \leq \ell_i$ for i moved earlier. $\ell'_j = f'_j - d_j = f_i - d_j \leq f_i - d_i \leq \ell_i$, so swapping inverted jobs does not increase max lateness. Thus there is some O with no inversions and no idle time by this. Since greedy has no idle time or inversions, it has the same lateness as O and is optimal. \square

Shortest Path

Goal: Find shortest directed path from s to other nodes.

```

1: procedure DIJKSTRA
2:    $S \leftarrow \{s\}$ 
3:    $\pi(s) = 0$ ,  $\pi(i) = \infty, i \neq s$ 
4:   Add nodes  $Q$  keys  $\pi(v)$ 
5:   while  $Q \neq \emptyset$  do
6:      $u = Q.ExtractMin()$ 
7:      $S \leftarrow S \cup \{u\}$ 
8:     for  $(u, v)$  with  $v \notin S$  and  $\pi(u) + \ell_{(u,v)} < \pi(v)$  do
9:        $Q.ChangeKey(v, \pi(v) = \pi(u) + \ell_{(u,v)})$ 

```

Proof. For each node $u \in S$, $\pi(u)$ is the shortest length of $s - u$. This is trivially true for $|S| = 1$. Suppose it is true for $|S| = k \geq 1$. Let v be the next node added to S and $u - v$ be the chosen edge. The shortest $s - u$ path plus (u, v) is an $s - v$ path of $\pi(v)$. Consider any other $s - v$ path P ; we will show it is not shorter than $\pi(v)$. Let $x - y$ be the first edge in P that leaves S , and let P' be the subpath to x . P is already too longer, for $\ell(P) \geq \ell(P') + \ell(x, y) \geq \pi(x) + \ell(x, y) \geq \pi(y) \geq \pi(x)$. \square

Minimum Spanning Tree

Def. A *path* is a sequences of edges which connects a sequence of nodes.

Def. A *cycle* is a path with no repeated nodes or edges besides the start and end nodes.

Def. An undirected graph is a *tree* if it is connected and does not contain a cycle.

Def. An undirected graph is a *spanning tree* if it is a tree and touches every vertex in G .

Goal: Given a connected graph G , find a tree T that is spanning with minimized sum of edge weights.

Prim

Greedly grow a tree from a root node.

Def. A *cut* is a partition of nodes into two nonempty sets S and $V - S$.

Def. The *cutset* of a cut S is the set of edges with exactly one endpoint in S .

Cut Property: Let S be any subset of nodes and let e be the min cost edge with exactly one endpoint in S . Then every MST contains e .

Proof. Given e from above and an MST T^* . Suppose $e \notin T^*$. Then $(u, v) = e$ must be connected by a path in T^* for it to be spanning. If we add e to T^* we create a cycle C . Thus $\exists f \in C, f \in \text{CutSet}(S)$. $T' = T^* \cup \{e\} - \{f\}$ is also a spanning tree and $c_e < c_f \Rightarrow \text{cost}(T') < \text{cost}(T^*)$, which is a contradiction. \square

```

1: procedure PRIM
2:   for  $v \in V$  do  $a[v] \leftarrow \infty$ 
3:    $Q$ 
4:   for  $v \in V$  do  $Q.add(v)$ 
5:    $S \leftarrow \emptyset$ 
6:   while  $Q \neq \emptyset$  do
7:      $u = Q.PullMin()$ 
8:      $S \leftarrow S \cup \{u\}$ 
9:     for  $e = (u, v)$  do
10:      if  $v \notin S$  &  $c_e < a[v]$  then
11:         $a[v] = c_e$ 

```

Reverse-Delete

Cycle property: Let C be any cycle, and let f be the max cost edge in C . Then $f \notin \text{MST}$.

Proof. Given $f \in C$ from above and an MST T^* . Suppose $f \in \text{MST}$. Deleting f disconnects T^* and creates a cut S . f is both in C and in $\text{CutSet}(S)$, so $\exists e \in C, e \in \text{CutSet}(S)$. $T' = T^* \cup \{e\} - \{f\}$ is also a spanning tree and $c_e < c_f \Rightarrow \text{cost}(T') < \text{cost}(T^*)$, which is a contradiction. \square

```

1: procedure REVERSE-DELETE
2:    $T \leftarrow E$ 
3:   Order  $e \in T$  in descending order of cost
4:   for  $e \in T$  do
5:     if  $T - \{e\}$  does not disconnect  $T$  then  $T - \{e\}$ 

```

Kruskal

```

1: procedure KRUSKAL
2:   Sort edge weights in descending order
3:    $T \leftarrow \emptyset$ 
4:   for  $u \in V$  do Create  $\{u\}$ 
5:   for  $i = 1$  to  $m$  do
6:      $(u, v) = e_i$ 
7:     if  $u, v$  in different sets then
8:        $T \leftarrow T \cup \{e_i\}$ 
9:       Merge sets containing  $u, v$ 
10:  return  $T$ 

```

Clustering

Goal: Given set U of n objects classify into coherent groups.

```

1: procedure CLUSTERING
2:   Form a graph on  $U$  with  $n$  clusters
3:   Find closest pair of objects in different clusters and add edge between them
4:   Repeat  $n - k$  times until there are  $k$  clusters

```

Divide-and-Conquer

Break up problems into several smaller parts and solve recursively.

```

1: procedure MS(L)
2:   Divide into halves  $L, R$ 
3:    $MS(L), MS(R)$ 
4:   Combine results

```

Recurrence:
 $T(n) = \begin{cases} 0 & n = 1 \\ T(\lceil n/2 \rceil) + T(\lfloor n/2 \rfloor) + n & \end{cases}$
If satisfied, $T(n) \leq n \lceil \lg n \rceil$.

Proof. Base case $n = 1$ is true. Let $n_1 = \lfloor n/2 \rfloor$ and $n_2 = \lceil n/2 \rceil$. Then $T(n) \leq T(n_1) + T(n_2) + n \leq n_1 \lceil \lg n_1 \rceil + n_2 \lceil \lg n_2 \rceil + n \leq n_1 \lceil \lg n_2 \rceil + n_2 \lceil \lg n_2 \rceil + n = n \lceil \lg n_2 \rceil + n \leq n(\lceil \lg n \rceil - 1) + n = n \lceil \lg n \rceil$. \square

Master Theorem

Let $T(n)$ be a monotonically increasing function with $T(n) = aT(n/b) + f(n)$, $T(1) = c$, where $a \geq 1$, $b \geq 2$, $c > 0$. Then if $f(n) \in \mathcal{O}(n^d)$ for $d \geq 0$:

$$T(n) = \begin{cases} \mathcal{O}(n^d) & \text{if } a < b^d \\ \mathcal{O}(n^d \log n) & \text{if } a = b^d \\ \mathcal{O}(n^{\log_b a}) & \text{if } a > b^d \end{cases}$$

This cannot be used if $T(n)$ is not monotone, if $f(n)$ is not polynomial, or if b cannot be expressed as a constant.

Dynamic Programming

Breaks problem into series of overlapping sub-problems and builds up a solution.

Coin-Changing

Goal: Make minimum change for S cents with a finite supply of coins from $C = \{v_1, \dots, v_n\}$. Solve $z(T, i)$ for minimum coins to reach $T \leq S$ with first i coins. Uses *memoization* to use values already computed.

```

1: procedure COINCHANGE
2:   for  $T = 1$  to  $S$  do
3:      $z(T, 0) \leftarrow \infty$ 
4:     for  $i = 0$  to  $n$  do
5:        $z(T, i) \leftarrow \infty$ 
6:       if  $T - v_i \geq 0$  then
7:          $z(T, i) \leftarrow \min(z(T, i), z(T - v_i, i))$ 

```

Weighted Interval Scheduling

Goal: find maximum weight subset of mutually compatible jobs.

Let $p(j)$ be the largest index $i < j$ such that i is compatible with j .

Then this has recurrence relation $OPT(j) = \begin{cases} 0 & j=0 \\ \max\{v_j + OPT(p(j)), OPT(j-1)\} & j>0 \end{cases}$

Bottom up algorithm:

```

1: procedure WINTSCH(L)
2:   Sort by increasing  $f_i$ 
3:   Compute  $p(i) \forall i$ 
4:    $M[0] = 0$ 
5:   for  $j = 1$  to  $n$  do
6:      $M[j] = \max(v_j + M[p(j)], M[j-1])$ 

```

Longest Common Subsequence

Goal: Given two strings, find the longest sequence of letters appearing in both, not necessarily contiguously.

$c[i, j] = \begin{cases} c[i-1, j-1] + 1 & x_i = y_j \\ \max\{c[i-1, j], c[i, j-1]\} & \text{otherwise} \end{cases}$

```

1: procedure LSC(X, Y)
2:    $m = \text{len}(X), n = \text{len}(Y)$ 
3:   for  $i = 1$  to  $m$  do  $c[i, 0] = 0$ 
4:   for  $j = 1$  to  $n$  do  $c[0, j] = 0$ 
5:   for  $i = 1$  to  $m$  do
6:     for  $j = 1$  to  $n$  do
7:       if  $x_i = y_j$  then
8:          $c[i, j] = c[i-1, j-1] + 1$ 
9:       else  $c[i, j] = \max\{c[i-1, j], c[i, j-1]\}$ 
10:  return  $c$ 

```

Sequence Alignment

Goal: Given two strings find alignment of minimum cost.

Def. An alignment M is a set of ordered pairs $x_i - y_j$ such that each item occurs in at most one pair and no crossings.

Def. The pair $x_i - y_j$ and $x_{i'} - y_{j'}$ cross if $i < i'$ but $j > j'$.

Let α be a mismatch cost and δ be an unmatched cost. Then:

```

1: procedure SEQA-
   LIGN(X, Y)
2:   for  $i = 0$  to  $m$  do
3:     for  $j = 0$  to  $n$  do
4:       for  $i = 1$  to  $m$  do
5:         for  $j = 1$  to  $n$  do
6:            $M[i, j] = \min(\alpha[x_i, y_j] + M[i-1, j-1], \delta + M[i, j-1])$ 
7:   return  $M[m, n]$ 

```

Knapsack

Goal: Fill knapsack to maximize total value.

```

1: procedure KNAPSACK
2:   for  $w = 0$  to  $W$  do
3:     for  $i = 1$  to  $n$  do
4:       for  $w = 1$  to  $W$  do
5:         if  $w_i > w$  then
6:            $M[i-1, w]$ 
7:         else
8:            $M[i, j] = \min(M[i-1, w], v_i + M[i-1, w-w_i])$ 
9:   return  $M[n, W]$ 

```

NPC

Class P consists of decision problems that are solvable in poly time. These problems are tractable. Otherwise problems are intractable.

Class NP consists of problems where a candidate solution can be verified in poly time. Y is NP-Hard if $X \leq_p Y \forall X \in NP$. Y is NPC if $Y \in NP$ and Y is NP-hard.

To establish Y is NPC:

1. Show $Y \in NP$
2. Choose X that is NPC and
3. Prove $X \leq_p Y$

Reduction from A to B is showing that we can solve A using the algorithm that solves B . Thus A is

easier than B . The idea is to transform an instance of A into inputs for B and show that they give identical answers.

NPC Problems

Minimum Vertex Cover: Find the smallest subset $S \subseteq V$ such that each edge has at least one endpoint in S .

Set Cover: Given a universe U and subsets S_i and integer k , does there exist a collection of S_i less than k such that $\bigcup S_i = U$.

3SAT: Given a conjunctive normal form Φ with three literals per clause, find an assignment of values to x_i so that Φ is True.

Independent Set: Given a graph, find the largest $S \subseteq V$ with no edges between them.

CLIQUE: Given a graph, is there a completely connected subgraph of size at least k .

HAM-CYCLE: Given an undirected graph, does there exist a simple cycle that contains every node in V .

TSP: Given a set of n cities and a pairwise distance function, is there a tour of length less than D .

Graph Coloring: Given a graph, can it be minimally colored so that every edge has different node colors.

Subset-Sum: Given natural numbers w_i and integer W , is there a subset of w_i that adds up to W

Solving NPC

Branch-and-Bound

Keeps track of best solution found so far (upper bound) and for each partial solution computes a lower bound on the objective function. If any bounds are not better than the best solution so far, they are pruned.

```

1: procedure BNB(P)
2:    $F \leftarrow \{(\emptyset, P)\}, B \leftarrow (\infty, (\emptyset, P))$ 
3:   while  $F \neq \emptyset$  do
4:     Choose  $(X, Y)$  in  $F$  and expand it to  $(X_i, Y_i)$  configurations
5:     for  $(X_i, Y_i)$  do
6:       if Solution
7:         found then
8:           if  $\text{cost}(X_i) < B$  then
9:              $B \leftarrow (\text{cost}(X_i), (X_i, Y_i))$ 
10:            if Not dead end then
11:              if  $LB(X_i) < B$  then
12:                 $F \leftarrow F \cup \{(X_i, Y_i)\}$ 
13:            return  $B$ 

```

Local Search

In this, the algorithm starts from an initial position then iteratively moves to a neighboring position using an evaluation function.

Hill-climbing: Choose neighbor with largest improvement as next state.

Stochastic: Randomize initialization step and search steps to allow for suboptimal choices.

Simulated Annealing: Select a neighbor at random. If it is better than current state, go there. Otherwise go there with some probability that decreases over time.

Tabu Search: In each step move to best neighbor solution even if worse than current. Avoids revisiting previously seen solutions with a tabu list of forbidden attributes.

Iterated LS: Generate an initial candidate then perform a LS on it. While termination conditions are not met, perturb current solution and do LS or return to original.

Approximation

In this, performance bounds are guaranteed with quick run-time. Def. The Ratio bound is

$$\max\left(\frac{A(x)}{OPT(x)}, \frac{OPT(x)}{A(x)}\right) \leq \rho(n)$$

for an input X of size n . Generally $\rho(n) \geq 1$, though unless $P = NP$ this equality will not hold.

For this, you can prove that an algorithm will give a solution no worse than $x * OPT$.

Integer Linear Programming

Given a_{ij}, b_i, c_i find x_j that satisfy $\min c'x$ subject to $Ax \geq b$ with x integral.

Network Flow

An abstraction for material flowing through edges. A directed graph without parallel edges with $c(e)$ equaling the capacity of edge e .

Def. An s - t flow is a function f from $E \rightarrow \mathbb{R}$ that satisfies

- For $e \in E: 0 \leq f(e) \leq c(e)$
- For $v \in V - \{s, t\}$: $\sum_{e \text{ into } v} f(e) = \sum_{e \text{ out of } v} f(e)$

Def. The value of the flow is $v(f) = \sum_{e \text{ out of } s} f(e)$.

Def. An s - t cut is a partition (A, B) of V with $s \in A$ and $t \in B$.

Def. The capacity of cut (A, B) is $cap(A, B) = \sum_{e \text{ out of } A} c(e)$.

Let f be any flow and let (A, B) be any s - t cut. Then $\sum_{e \text{ out of } A} f(e) = v(f)$.

Weak duality: Let f be any flow and let (A, B) be any s - t cut. Then the value of the flow is at most the capacity of the cut. Corollary: If $v(f) = cap(A, B)$, then f is a max flow and (A, B) is a min s - t cut.

Def. A residual edge of $e = (u, v)$ is $e^R = (v, u)$ with $c_f(e) = \begin{cases} c(e) - f(e) & e \in E \\ f(e) & e^R \in E \end{cases}$

1: procedure FORD-
FULKERSON
2: for $e \in E$ do $f(e) \leftarrow 0$
3: $G_f \leftarrow$ residual graph
4: while $\exists s-t$ path $P \in G_f$ do
5: $f \leftarrow Aug(f, c, P)$
6: update G_f
7: return f
8: procedure AUG(f, c, P)
9: $b \leftarrow$ bottleneck(P)
10: for $e \in P$ do
11: if $e \in E$ then $f(e) \leftarrow f(e) + b$
12: else $[e^R \in E] f(e^R) \leftarrow f(e^R) - b$
13: return f

Augmenting path theorem: Flow f is a max flow \iff there are no augmenting paths.

Max-flow min-cut theorem: The value of the max flow is equal to the value of the min s - t cut. We have equivalence between

1. $\exists(A, B)$ such that $v(f) = cap(A, B)$
2. Flow f is a max flow
3. There is no augmenting path relative to f

Bipartite Matching

Given an undirected, bipartite graph, find the max cardinality matching.

Can be reduced to a max flow problem by adding nodes s, t with edge capacities 1.

Theorem: Max cardinality matching equals value of max flow.

Disjoint Paths

Given a directed graph and nodes s, t , find the max number of edge-disjoint (no edge in common) s - t paths.

Theorem: Max number edge-disjoint s - t paths equals max flow by assigning all $c(e) = 1$.

Network Connectivity

Given a digraph and nodes s, t , find min number of edges whose removal disconnects s - t .

Theorem: The max number of edge-disjoint s - t paths is equal to the min number of edges whose removal disconnects s - t .

Choosing path with highest bottleneck capacity increases flow by max possible amount.

```

1: procedure MAX-FLOW
2:   for  $e \in E$  do  $f(e) \leftarrow 0$ 
3:    $\Delta \leftarrow \min 2^n \geq c$ 
4:    $G_f \leftarrow$  residual graph
5:   while  $\Delta \geq 1$  do
6:      $G_f(\Delta) \leftarrow \Delta - G_{res}$ 
7:     while  $\exists P \in G_f(\Delta)$  do
8:        $f \leftarrow aug(f, c, P)$ 
9:       update  $G_f(\Delta)$ 
10:     $\Delta \leftarrow \Delta/2$ 
11:  return  $f$ 

```