

CSE 6140/ CX 4140:

Computational Science and Engineering

ALGORITHMS

Instructor: Anne Benoit


Visiting Associate Professor, CSE

Based on slides by Bistra Dilkina

KT4.5 Minimum Spanning Trees



Slides by Kevin Wayne.
Copyright © 2005 Pearson-Addison Wesley.
All rights reserved.



Trees

Def. A **path** is a sequence edges which connects a sequence of nodes.


Def. A **cycle** is a path with no repeated nodes or edges other than the start and end nodes.

Def. An undirected graph is a **tree** if it is connected and does not contain a cycle.

Def. An undirected graph is a **spanning tree** if it is a tree and touches every vertex in G .


Theorem. Let G be an undirected graph on n nodes. Any two of the following statements imply the third.

- G is connected.
- G does not contain a cycle.
- G has $n-1$ edges.




Minimum Spanning Tree

Minimum spanning tree. Given a connected graph $G = (V, E)$ with real-valued edge weights c_e , an MST is a subset of the edges $T \subseteq E$ such that T is a spanning tree whose sum of edge weights is minimized.



$G = (V, E)$



$T, \sum_{e \in T} c_e = 50$

Cayley's Theorem. There are n^{n-2} spanning trees of K_n .

↑
can't solve by brute force

Applications

MST is fundamental problem with diverse applications.

- Network design.
 - telephone, electrical, hydraulic, TV cable, computer, road
- Approximation algorithms for NP-hard problems.
 - traveling salesperson problem, Steiner tree
- Indirect applications.
 - max bottleneck paths
 - learning salient features for real-time face verification
 - reducing data storage in sequencing amino acids in a protein
 - model locality of particle interactions in turbulent fluid flows
- Cluster analysis.

Greedy Algorithms

Kruskal's algorithm. Start with $T = \emptyset$. Consider edges in ascending order of cost. Insert edge e in T unless doing so would create a cycle.

Reverse-Delete algorithm. Start with $T = E$. Consider edges in descending order of cost. Delete edge e from T unless doing so would disconnect T .

Prim's algorithm. Start with some root node s and greedily grow a tree T from s outward. At each step, add the cheapest edge e to T that has exactly one endpoint in T .

Remark. All three algorithms produce an MST.


Greedy Algorithms

Def. A **cut** is a partition of the nodes into two nonempty subsets S and $V - S$.


Def. The **cutset** of a cut S is the set of edges with exactly one endpoint in S .

Simplifying assumption. All edge costs c_e are distinct.

Cut property. Let S be any subset of nodes, and let e be the min cost edge with exactly one endpoint in S . Then every MST contains e .



Cut $S = \{4, 5, 8\}$
Cutset $D = 5-6, 5-7, 3-4, 3-5, 7-8$



e is in the MST


Cut Property

Cut property. Let S be **any** subset of nodes, and let $e=(u,v)$ be the min cost edge with exactly one endpoint in S . Then every MST T^* contains e .

Simplifying assumption. All edge costs c_e are distinct.

Pf. (exchange argument)


- Given: e is the min cost edge in the cutset for a set S , T^* is an MST
- Suppose e does not belong to T^* , and let's see what happens.
- The endpoints u and v of e must be connected by a path in T^* (spanning)
- \Rightarrow Adding e to T^* creates a cycle C in T^*
- Edge e is both in the cycle C and in the cutset corresponding to $S \Rightarrow$ there exists another edge, say f , that is in both C and cutset of S .
- $T' = T^* \cup \{e\} - \{f\}$ is also a spanning tree, because:
 - All n nodes are connected (f, e were part of a cycle) AND $n-1$ edges
- Since $c_e < c_f$, $\text{cost}(T') < \text{cost}(T^*)$.
- This is a contradiction since T^* is an optimal spanning tree, our assumption about e not being in T^* must be wrong.



Prim's Algorithm

Prim's algorithm. [Jarník 1930, Dijkstra 1957, Prim 1959]

- Initialize S = any node.
- (Apply cut property to S .)
- Add to tree the **min cost edge (u,v) in cutset** corresponding to S , and add one new explored node u to S .




Prim's Algorithm: Proof of Correctness

Prim's algorithm. [Jarník 1930, Dijkstra 1957, Prim 1959]

- Initialize $S = \text{any node}$.
- (Apply cut property to S .)
- Add to tree the min cost edge (u,v) in cutset corresponding to S , and add one new explored node u to S .

1) Every edge added satisfies the Cut Property: by design

2) It produces a spanning tree: alg stops when $S=V$



Implementation: Prim's Algorithm

Implementation. Use a priority queue (as for Dijkstra).

- Maintain set of explored nodes S .
- For each unexplored node v , maintain attachment cost $a[v]$ = cost of cheapest edge v to a node in S .
- $O(n^2)$ with an array; $O(m \log n)$ with a binary heap.

```
Prim(G, c) {  
  foreach (v ∈ V) a[v] ← ∞  
  Initialize an empty priority queue Q  
  foreach (v ∈ V) insert v onto Q  
  Initialize set of explored nodes S ← ∅  
  
  while (Q is not empty) {  
    u ← delete min element from Q  
    S ← S ∪ { u }  
    foreach (edge e = (u, v) incident to u)  
      if ((v ∉ S) and (ce < a[v]))  
        decrease priority a[v] to ce  
  }
```

Shortest edge between v and
a node in explored set S




Greedy Algorithms


Simplifying assumption. All edge costs c_e are distinct.

Cycle. Set of edges the form $a-b, b-c, c-d, \dots, y-z, z-a$.

Cycle property. Let C be any cycle, and let f be the max cost edge belonging to C . Then every MST does not contain f .



Cycle $C = 1-2, 2-3, 3-4, 4-5, 5-6, 6-1$



f is not in the MST


Cycle Property

Cycle property. Let C be any cycle in G , and let f be the max cost edge belonging to C . Then every MST T^* does not contain f .

Simplifying assumption. All edge costs c_e are distinct.

Pf. (exchange argument)

- Given: f is the max cost edge in a cycle C , T^* is an MST
- Suppose f belongs to T^* , and let's see what happens.
- Deleting f from T^* disconnects T^* and creates a cut S in T^* .
- Edge f is both in the cycle C and in the cutset D corresponding to S
 \Rightarrow there exists another edge, say e , that is in both C and D .
- $T' = T^* \cup \{e\} - \{f\}$ is also a spanning tree (same argument as before).
- Since $c_e < c_f$, $\text{cost}(T') < \text{cost}(T^*)$.
- This is a contradiction since T^* is an optimal spanning tree, hence our assumption about f is wrong.



Reverse-Delete Algorithm: Proof of Correctness


The algorithm. Start with $T = E$. Consider edges in descending order of cost. Delete edge e from T unless doing so would disconnect T .

Th. This algorithm produces an MST of G .

Cycle property. Let C be any cycle in G , and let f be the max cost edge belonging to C . Then every MST T^* does not contain f .

Pf.


- Edges removed do not belong to any MST.
- The output is a spanning tree: connected & no cycle.




Kruskal's Algorithm: Proof of Correctness

Kruskal's algorithm. [Kruskal, 1956]

- Consider edges in ascending order of weight.
- **Case 1:** If adding e to T creates a cycle, discard e (according to cycle property: because of sorted order, e must be maxcost in cycle).
- **Case 2:** Otherwise, insert $e = (u, v)$ into T (according to cut property: for set $S =$ set of nodes in u 's connected component in current set T).



Case 1



Case 2

Implementation: Kruskal's Algorithm

Implementation. Use the **union-find** data structure.

- Build set T of edges in the MST.
- Maintain a set for each connected component.
- $O(m \log n)$ for sorting and $O(m \underbrace{\alpha(m, n)}_{\text{essentially a constant}})$ for union-find.

$m \leq n^2 \Rightarrow \log m$ is $O(\log n)$

essentially a constant

```
Kruskal(G, c) {  
    Sort edges weights so that  $c_1 \leq c_2 \leq \dots \leq c_m$ .  
     $T \leftarrow \phi$   
  
    foreach ( $u \in V$ ) make a set containing singleton  $u$   
  
    for  $i = 1$  to  $m$     are  $u$  and  $v$  in different connected components?  
        ( $u, v$ ) =  $e_i$     ↙  
        if ( $u$  and  $v$  are in different sets) {  
             $T \leftarrow T \cup \{e_i\}$   
            merge the sets containing  $u$  and  $v$   
        }  
    ↙ merge two components  
    return  $T$   
}
```


Lexicographic Tiebreaking

Simplifying assumption. All edge costs c_e are distinct.

Impact. Kruskal and Prim only interact with costs via pairwise comparisons.

To remove the assumption that all edge costs are distinct: perturb all edge costs by tiny amounts to break any ties.

Implementation. Can handle arbitrarily small perturbations implicitly by breaking ties lexicographically, according to index.

```
boolean less(i, j) {  
    if      (cost(ei) < cost(ej)) return true  
    else if (cost(ei) > cost(ej)) return false  
    else if (i < j)                 return true  
    else                           return false  
}
```