

CSE 6140/ CX 4140:


Computational Science and Engineering

ALGORITHMS

Instructor: Anne Benoit
Visiting Associate Professor, CSE

Based on slides by Bistra Dilkina and
Holger Hoos

HEURISTICS/LOCAL SEARCH [SLS2]



'This material is based on slides provided with the book 'Stochastic Local Search: Foundations and Applications' by Holger H. Hoos and Thomas Stützle (Morgan Kaufmann, 2004) - see www.sls-book.net for further information.'


Local Search (LS) Algorithms

- search space S
 - SAT: set of all complete truth assignments to propositional variables (all “potential solutions”)
- solution set $S' \subseteq S$
 - SAT: all satisfying assignments for given formula
- neighborhood relation $N \subseteq S \times S$
 - A way to move from one potential solution to another
 - SAT: neighboring variable assignments differ in the truth value of exactly one variable
- evaluation function $g : S \rightarrow \mathbb{R}^+$
 - SAT: number of clauses satisfied under given assignment


Local Search (LS)

- Start from initial position
- Iteratively move from current position to one of neighboring positions
- Use evaluation function to choose among neighboring positions

Local Search




Global view of search space



Local view of search space

State space landscape

- Objective function defines *state space landscape*




Local Search (LS) Algorithms

- search space S
 - SAT: set of all complete truth assignments to propositional variables (all “potential solutions”)
- solution set $S' \subseteq S$
 - SAT: all satisfying assignments for given formula
- neighborhood relation $N \subseteq S \times S$
 - A way to move from one potential solution to another
 - SAT: neighboring variable assignments differ in the truth value of exactly one variable
- evaluation function $g : S \rightarrow \mathbb{R}^+$
 - SAT: number of clauses satisfied under given assignment

Local Search (LS) Algorithms

- search space S
 - TSP: set of all permutations of vertices (all “potential solutions”)
- solution set $S' \subseteq S$
 - TSP: the tours of minimum length
- neighborhood relation $N \subseteq S \times S$
 - A way to move from one potential solution to another
 - TSP: neighboring tour differ in several edges
- evaluation function $g : S \rightarrow \mathbb{R}^+$
 - TSP: length of tour

Symm. TSP --- search neighborhood



Search Space: all permutations of the cities (each defines a cycle)

3-opt – delete 3 edges and reconnect fragments into 1 cycle


k-opt – delete k edges and reconnect fragments into 1 cycle

Iterative Improvement (Greedy Search)

- Initialize search at some point of search space
- At each step, move from the current search position to a neighboring position with better evaluation function value

Hill climbing (Best Improvement Search)

- Choose the neighbor with the largest improvement as the next state



```
while f-value(state) < f-value(next-best(state))  
    state := next-best(state)
```

function Hill-Climbing(*problem*) **returns** a *solution state*

current \leftarrow Make-Node(Initial-State[*problem*])

loop do

next \leftarrow a highest-valued successor of *current*

if Value[*next*] < Value[*current*] **then return** *current*

current \leftarrow *next*

end

Problems with iterative improvement

- Advantages:
 - Very fast, works well for certain problems
- Disadvantages:
 - What if there are multiple peaks?
 - Hill climbing gets stuck at all peaks, known as local maxima
 - Optimal solution is highest peak – global maximum
 - May result in extremely suboptimal solution if many peaks
 - Being misguided by evaluation/objective function

Stochastic Local Search

- randomize initialization step
- randomize search steps such that suboptimal/worsening steps are allowed
- improved performance & robustness
- typically, degree of randomization controlled by noise parameter

Stochastic Local Search


- Pros:
 - for many combinatorial problems, more efficient than systematic search
 - easy to implement
 - easy to parallelize
- Cons:
 - often incomplete (no guarantees for finding existing solutions)
 - highly stochastic behavior
 - often difficult to analyze theoretically/empirically

Simple SLS methods

- Random Search (Blind Guessing):
 - In each step, randomly select one element of the search space.
- (Uninformed) Random Walk:
 - In each step, randomly select one of the neighboring positions of the search space and move there.

Random restart hill climbing

- Start at random solution
- Hill-climb until local optima
- Start at another random position



Randomized Iterative Improvement:

- Idea: escape local maxima by allowing some "bad" moves
- initialize search at some point of search space
- search steps:
 - with probability p , move from current search position to a randomly selected neighboring position
 - otherwise, move from current search position to neighboring position with better evaluation function value
- Has many variations of how to choose the random neighbor, and how many of them

The WalkSAT Algorithm

```
WalkSAT(CNF, max-tries, max-flips, p) {  
    for i ← 1 to max-tries do  
        solution = random truth assignment  
        for j ← 1 to max-flips do  
            if all clauses in CNF satisfied then  
                return solution  
            c ← random unsatisfied clause in CNF  
            with probability p  
                flip a random variable in c  
            else  
                flip variable in c that maximizes  
                    number of satisfied clauses  
        return failure  
}
```

Simulated annealing (SA)

- Combinatorial search technique inspired by the physical process of annealing [Kirkpatrick et al. 1983, Cerny 1985]
- Outline
 - Select a neighbor at random
 - If better than current state, go there (improving move)
 - Otherwise, go there with some probability (worsening move)
 - Probability goes down with time (similar to temperature cooling)
- When probability is high → diversify (many worsening moves)
- When probability is low → intensify (focus on improving moves)

SA analogy

- Annealing is process of heating and cooling metals in order to improve strength
- Idea: Controlled heating and cooling of metal
 - When hot, atoms move around
 - When cooled, atoms find configuration with lower internal energy (i.e. makes metal stronger)
- Analogy:
 - Temperature = probability of accepting worse neighboring solution
 - When temperature is high, likely to accept worse neighboring solutions (but may lead to better overall solution)
 - Analogous to atoms wandering around
 - Cooling represents shrinking probability of accepting worse solutions

SA Pseudo code

function Simulated-Annealing(*problem, schedule*) **returns**
solution state

current \leftarrow Make-Node(Initial-State[*problem*])

for *t* \leftarrow 1 **to** *infinity* (*Iters, Time cutoff*)

T \leftarrow *schedule*[*t*] // *T* goes downwards.

if *T* = 0 **then** **return** *greedy from current*

next \leftarrow Random-Successor(*current*)

ΔE \leftarrow f-Value[*next*] - f-Value[*current*]

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* with probability $e^{\Delta E/T}$

end

Simulated annealing (SA)

- **Acceptance criterion (Metropolis condition):** choose new solution s' over old solution s with probability (maximization)

$$\Pr(s', s) = \begin{cases} 1 & \text{if } f(s') > f(s) \\ \exp\left\{\frac{f(s') - f(s)}{T}\right\} & \text{otherwise} \end{cases}$$

- **Initial temperature** T_0
- **Annealing (cooling) schedule:** how to update the temperature
 - E.g. $T = a T$ with $a = 0.95$ (geometric schedule)
 - Number of iterations at each temperature (e.g. multiple of the neighborhood size)
- **Stopping criterion**
 - E.g. no improved solution found for a number of iterations (or number of temperature values)

SA for TSP [Johnson & McGeoch 1997]

- baseline implementation:
 - start with random initial solution
 - use 2-exchange neighborhood
 - simple annealing schedule
- → relatively poor performance
- improvements:
 - look-up table for acceptance probabilities
 - neighborhood pruning
 - low-temperature starts

SA with restarts

- RESTARTS: Sometimes it is better to move back to a solution that was significantly better rather than always moving from the current state.
- The decision to restart could be based on several criteria.
 - based on a fixed number of steps,
 - based on whether the current energy is too high compared to the best energy obtained so far,
 - too many iterations without improvement,
 - restarting randomly, etc.

Summary of Simulated Annealing

- is historically important
- is easy to implement
- has interesting theoretical properties (convergence), but these are of very limited practical relevance
- achieves good performance often at the cost of substantial run-times

Tabu Search

- Combinatorial search technique that heavily relies on the use of an explicit memory of the search process [Glover 1989, 1990] to guide search process
- memory typically contains only specific attributes of previously seen solutions
- simple tabu search strategies exploit only short term memory
- more complex tabu search strategies exploit long term memory

Tabu search – exploiting short term memory

- in each step, move to best neighboring solution although it may be worse than current one
- to avoid cycles, tabu search tries to avoid revisiting previously seen solutions by basing the memory on attributes of recently seen solutions
- tabu list stores attributes of the TL most recently visited solutions; parameter TL is called tabu list length or tabu tenure
- solutions that contain tabu attributes are forbidden

Tabu Search

- Problem: previously unseen solutions may be tabu → use of **aspiration criteria** to override tabu status
- Stopping criteria:
 - all neighboring solutions are tabu
 - maximum number of iterations exceeded
 - number of iterations without improvement

Example: Tabu Search for SAT / MAX-SAT

- Neighborhood: assignments that differ in exactly one variable instantiation
- Tabu attributes: variables
- Tabu criterion: flipping a variable is forbidden for a given number of iterations
- Aspiration criterion: if flipping a tabu variable leads to a better solution, the variable's tabu status is overridden
- [Hansen & Jaumard 1990; Selman & Kautz 1994]

Iterated local search


- Generate initial candidate solution s
- Perform local search on s (for example iterative improvement starting from s)
- While termination condition not met
 - Set $r=s$
 - Perform perturbation on s
 - Perform local search on perturbed s
 - Based on acceptance criterion, keep s or revert to r

Iterated local search

- ILS can be interpreted as walks in the space of local optima
- Perturbation is key
 - Needs to be chosen so that it cannot be undone easily by subsequent local search
 - It may consist of many perturbation steps
 - Strong perturbation: more effective escape from local optima but similar drawbacks as random restart
 - Weak perturbation: short subsequent local search phase but risk of revisiting previous optima
- Acceptance criteria: usually either the most recent or the better of two

Iterated local search for TSP

- Perturbation: “double-bridge move” = 4-exchange step
- Cannot be directly reversed by 2-exchange moves



Other search techniques

- Genetic algorithms
- Ant colony optimization
- Usually covered in AI courses

Construction heuristics for initial solutions

- search space: space of partial solutions
- search steps: extend partial solutions with assignment for the next element
- solution elements are often ranked according to a greedy evaluation function

TSP construction: Nearest neighbor

- Start at some vertex s ; $v=s$;
- While not all vertices visited
 - Select closest unvisited neighbor w of v
 - Go from v to w
 - $v=w$
- Go from v to s
- Running time $O(n^2)$

TSP construction: Many variants

- **Closest insertion:** insert vertex closest to vertex in the tour
- **Farthest insertion:** insert vertex whose minimum distance to a node on the cycle is maximum
- **Cheapest insertion:** insert the node that can be inserted with minimum increase in cost
- **Random insertion:** randomly select a vertex and insert vertex at position that gives minimum increase of tour length

CSE 6140/ CX 4140

Empirical Analysis of Algorithms

[SLS4]

textbook: STOCHASTIC LOCAL SEARCH
FOUNDATIONS AND APPLICATIONS

based on slides by Holger Hoos

Theoretical vs. Empirical Analysis

Ideal: Analytically prove properties of a given algorithm
(run-time: worst-case / average-case / distribution, error rates).

Reality: Often only possible under substantial simplifications or not at all.

~> Empirical analysis

The Three Pillars of CS:

- **Theory:** abstract models and their properties
("eternal truths")
- **Engineering:** principled design of artifacts
(hardware, systems, algorithms, interfaces)
- **(Empirical) Science:** principled study of phenomenae
(behaviour of hardware, systems, algorithms; interactions)

The Scientific Method

make observations

formulate hypothesis/hypotheses (model)

While not satisfied (and deadline not exceeded) iterate:

1. design experiment to falsify model
2. conduct experiment
3. analyse experimental results
4. revise model based on results

Goals

- Defining standard methodologies
- Comparing relative performance of algorithms so as to identify the best ones for a given application
- Characterizing the behavior of algorithms
- Identifying algorithm separators, *i.e.*, families of problem instances for which the performance differ
- Providing new insights in algorithm design

Issues:

- algorithm implementation (fairness)
- selection of problem instances (benchmarks)
- performance criteria (what is measured?)
- experimental protocol
- data analysis & interpretation

Benchmark Selection

Some criteria for constructing/selecting benchmark sets:

- instance hardness (focus on hard instances)
- instance size (provide range, for scaling studies)
- instance type (provide variety):
 - individual application instances
 - hand-crafted instances (realistic, artificial)
 - ensembles of instances from random distributions
(\leadsto random instance generators)
 - encodings of various other types of problems
(*e.g.*, SAT-encodings of graph colouring problems)

CPU Time *vs.* Elementary Operations

How to measure run-time?

- Measure CPU time (using OS book-keeping & functions)
- Measure elementary operations of algorithm
(*e.g.*, local search steps, calls of expensive functions)
and report cost model (CPU time / elementary operation)

Issues:

- accuracy of measurement
- dependence on run-time environment
- fairness of comparison